

The Modern AI-Discoverable SaaS Launch Playbook

(INTRO)

The condensed intro — five pages on the 5-layer model + three named defensive patterns. The full 100+ page playbook is at hyder.me/playbook.

Foreword: Why this playbook exists

Building for the web in 2026 is fundamentally different from building for the web in 2022.

In 2022, "discoverability" meant one thing: Google. You wrote keyword-targeted content, you earned backlinks, you waited for rankings to compound. Bing, Yandex, DuckDuckGo existed but were rounding errors.

That world is gone. The traffic you used to get from a #1 Google ranking now splits across at least seven distinct surfaces:

1. **Google web search** — still ~60% of organic traffic, but increasingly stripped to a snippet by Google AI Overviews
2. **Google AI Overviews** — pulls from a different stack than blue-link search; lives or dies by structured data
3. **Perplexity** — grounded answer engine; cites sources visibly; trained on the public web with heavy crawl
4. **ChatGPT search** — same family of behaviors as Perplexity, different surfacing logic
5. **Claude (Anthropic) web tools** — grounding from documents the user attaches, real-time fetches, and the model's training corpus
6. **Bing/Copilot** — small organic share, but where IndexNow actually lives + where Bing AI grounds
7. **Vertical AI tools** — every domain-specific AI (code agents, research bots, finance tools, etc.) is increasingly using OpenAI tool-calling against APIs you might not even know about

Each of these surfaces consumes a different signal. Google AI Overview wants `FAQPage` `JSON-LD`. Perplexity wants clean canonical text and Wikidata grounding. Claude's tool calling wants a polished OpenAPI 3.1 spec. Bing wants IndexNow pings.

The new mandate: **build for every surface simultaneously, with one set of files**. That's what this playbook is.

It's organized in five layers (plus a pre-flight chapter), built from foundation up. Skip nothing — they compound.

§1. The 5-layer model (TLDR)

Layer 5: Distribution Press, social, journalist outreach, directory listings
Layer 4: Conversion + analytics GA4 events, success URLs, error handling
Layer 3: Compliance + headers CSP, HSTS, COOP, Permissions-Policy, RLS, secrets
Layer 2: Search + answer engines JSON-LD, OpenAPI, sitemaps, IndexNow, FAQs
Layer 1: Entity identity Wikidata, llms.txt, canonical IDs, structured grounding

Each layer requires the one below it. You can't get good GA4 conversion attribution if your URLs aren't canonical (Layer 2). You can't get into Google AI Overviews if your headers leak mixed content (Layer 3). You can't pitch journalists effectively if your Wikidata entry is broken (Layer 1).

Build bottom-up. Reading top-down (skipping foundations) is the most common reason launches don't compound.

A taste of §8 — defensive patterns

The full playbook has 12 named defensive patterns from real production bugs. Each one is organized by symptom so you can grep the playbook when something breaks in your codebase. Three of the most-painful:

8.1 The Stripe newline bug

Symptom: SDK calls fail with "Connection error, retried 2 times" or "no such price: 'price_XXX\n'". **Root cause:** Env var has trailing newline from `echo` piping. **Fix:** Use `printf %s` not `echo` when setting; `.trim()` every read in code.

8.3 The gtag race condition

Symptom: Conversion events fire (according to your code) but never appear in GA4 Realtime. **Root cause:** `window.gtag` is undefined when your `useEffect` runs. **Fix:** Push to `window.dataLayer` directly. The `dataLayer` is initialized synchronously by the GA loader; the `gtag` function is hydrated later.

8.5 The slug collision 404

Symptom: URLs return 404 even though the resource exists. **Root cause:** Multiple resources have the same slug (e.g., the same matchup repeats across days, daily snapshots, recurring meetings). `maybeSingle()` returns null when multiple rows match. **Fix:** Append a uniqueness suffix to the slug for repeating resource types (e.g., date for sports games, instance ID for recurring events).

§11. Closing thoughts

The web is changing under our feet faster than at any point since mobile.

LLMs are the new search engines. Wikidata is the new Encyclopedia Britannica. JSON-LD is the new sitemap.xml. OpenAPI is the new API documentation page. Hugging Face is the new data partnerships page. Direct journalist outreach is the new SEO.

What hasn't changed: building a thing people want, telling people about it clearly, and getting out of their way. Layer 5 (distribution) without Layers 1-4 (foundation) is shouting in an empty room. Layers 1-4 without Layer 5 is a library nobody visits.

Build all five. The compounding is what works.

Get the full playbook

This is 5 pages of a 100+ page playbook. The full version covers:

- **§0 Pre-flight** — domain, DNS, GA4, Search Console, Bing WMT, SPF/DKIM/DMARC
- **Layers 1-5** — Wikidata, JSON-LD, OpenAPI, security headers, GA4 events, distribution
- **§8** — all 12 named defensive patterns (you've seen 3)
- **§9** — pre-launch + weekly + recovery checklists
- **Templates** — llms.txt, JSON-LD generator, reference CSP, journalist pitches
- **Claude skill** — apply the playbook to any project automatically
- **12 months of free updates**

Available at hyder.me/playbook for \$79. 7-day refund.

If you'd rather have Kenny apply the playbook to your project directly, the done-for-you launch engagement is at hyder.me/contact?topic=playbook-dfy.

Copyright © 2026 Kenny Hyder · Hyder Media · hyder.me

This playbook is published by Hyder Media. Personal-use license within one organization.
For distribution rights, email kenny@hyder.me.